



So, at any given time, memory (representing the true database state) and files will not exactly contain the same data.



What happens if the system crashes and everything that was in memory simply vanishes?



39

Not committed?

If the change wasn't committed, then losing it is just what we would expect: a crash should rollback any on-going transaction.

Committed?


If the change was committed it's a different story. The D in ACID means that we expect any committed change to survive a system crash.

DURABILITY

40


Practically, it means that every COMMIT should result into a SYNCHRONOUS file write: the commit call only returns when changes are safely written to disk. We aren't going to update all the blocks affected by the transaction, which would take too much time, but only save the new values to an additional log (journal) file.

commit




41

In Oracle, every pending change (new values, as opposed to old values kept for rollback) will also be saved to a "log buffer" that will be flushed to a log file on commit (or when dangerously full)




When you shutdown the database cleanly, all the buffers are flushed to files, which represent the last state in memory.

Clean shutdown :
Flush buffers to files



When restarting after a crash, the log is scanned for all committed changes, and they are applied if need be to the file blocks, until the files contain the exact image of the database at the last COMMIT before the crash.

RESTART AFTER CRASH



44

What about several sessions?

Dedicated Sessions

Assign to all connected users their own personal server process running SQL for them.

Session Pooling

Pre-start (which makes connection slightly faster) a number of servers, and to assign requests to them with a bit of load-balancing.

What about a session querying data being changed by another session?

The main problem is that a transaction can last "a certain time". You may need to run a lot of code before deciding on commit or rollback.

```
Begin Transaction
Commit
```

During the transaction, the database will be in a kind of transient state of which you cannot say whether it will become permanent.

CONSISTENT STATE

Between two consistent states, you have a kind of black hole.

4 levels for ISOLATION

The SQL standard defines four isolation levels, from no isolation at all to paranoid. Some products let you set it, others impose it.

- ~~dirty reads~~
- read committed
- repeatable read
- serialization

Let's see once again how a transaction works. You start it. Before applying any change the DBMS saves the value from before your transaction, in case you'd want to rollback. Then you commit (or rollback) and the value previously saved for rollback (in memory or on file) can be marked as disposable.

~~Read Uncommitted (dirty read)~~

If when a user is changing data another user were able to read the new value before being sure that there will be no rollback, we would run into problems pretty soon.

Read Committed

But the "old" value belongs to a stable (committed) state of the database. What most DBMS products will do is that they will "serve" to the reader this value, at least known to be consistent and to have been once "official current value".

Repeatable read

Point-in-time read

However, once the writer has committed, in the read committed model whe should read the new, now official, value. It may be inconsistent with a previous read and we may favor consistency for the reader over timeliness.

In practice, the problem is more a problem of data consistency of foreign keys when we are scanning big related tables. A single SELECT will usually be consistent (a product such as Oracle ignores any change having happened since the start of the SELECT, even if it was committed). But if we SELECT twice (two different queries) from two tables with an FK relationship, changes that may have happened (and have been committed) between the time when we started reading the first table and the time when we were reading the second table may lead to problems such as orphaned rows.

Example: If I haven't read an uncommitted order when reading ORDER, then I should ignore rows refering to this order in ORDER_DETAIL even if they have been committed by the time I reach them.

Serialization

The last level is to block readers when data is being modified. Guaranteed to be consistent but bad for performances.

When you have twowritersyou cannot do otherwise than locking.

BACKUP ISSUES

SCALING UP

Adding servers

Neil Gunther's Universal Law of Computational Scalability

Lab 11

[Slides](#)

[INFORMATION_SCHEMA \(PostgreSQL Version\)](#)

Data Dictionary / Catalog

A set of tables that contain information about the objects in the database.

One catalog per database.

System tables

Any database stores "metadata" that describes the tables in your database (and not only them).

All client tools use this information to let you browse the structure of your tables

Whenever you are issuing DDL commands, you are actually modifying system tables. They must NEVER be directly changed.

Read access to these tables is provided through system views.

In these views you only see what YOU are allowed to see. Only administrators see everything.

Most important tables for developers in INFORMATION_SCHEMA (PostgreSQL version)

- TABLES
- COLUMNS
- TABLE_CONSTRAINTS
- CONSTRAINT_COLUMN_USAGE
- KEY_COLUMN_USAGE
- CHECK_CONSTRAINTS

SEQUENCES ROUTINES TRIGGERS VIEWS

VIEW_TABLE_USAGE

Those are all views you should be aware of.

Nothing on indexes not associated with PK or UNIQUE constraint!

Must look into `pg_catalog`.

- `pg_index`
- `pg_class`
- `pg_attribute`

Stuff mostly for DB administrators

Roles and privileges

- `pg_class` relates tables to files on disk
- `pg_statistic` estimates used by the optimizer
- `pg_settings` database parameters

...